

**STAT 453**

**SS 2024**

**Midterm Exam**

**03/19/2024**

**Time:** 4:00 – 5:15 pm am (75 mins) **Email:** \_\_\_\_\_ @wisc.edu

**Instructor:** Yiqiao Zhong

**Teaching Assistant:** Zhexuan Liu

**Access Code:** 819771-ZHONG

---

**Name:** \_\_\_\_\_

This exam contains 7 pages (including this cover page) and 9 questions.

Total of points is 100.

By submitting this exam, I (the student)

- acknowledge that I am required to follow the academic integrity and conduct policies of UW-Madison.

Grade Table (for teacher use only)

Question	Points	Score
1	10	
2	10	
3	10	
4	10	
5	10	
6	10	
7	10	
8	10	
9	20	
Total:	100	

---

1. (10 points) Among the four logical gates, which one cannot achieve 100% using a linear decision boundary?

- A. AND
  - B. OR
  - C. NOT
  - D. XOR
- 

Solution:

2. (10 points) Which statement is wrong about stochastic gradient descent (SGD) for training neural networks?

- A. Compared with GD, SGD is much faster due to fewer examples a network computes at each iteration.
  - B. Due to noise, SGD shows less oscillatory behavior compared with GD.
  - C. SGD and perceptron algorithm share many similarities.
  - D. When training neural networks, the standard practice is to only use first-order derivatives for SGD.
- 

Solution:

3. (10 points) CNNs are more effective than MLPs in processing images. Which one is not a reason for the effectiveness of CNNs?

- A. CNNs benefit from hand-crafted features to extract features such as edges corners from images.
- B. The kernels of CNNs use patches to compute local information, which reduces the number of parameters.
- C. CNNs have kernels that are shared across all patches, which benefit from our prior knowledge that images are invariant to translations.
- D. Many network-in-network modules such as those in Inception further reduce the number of parameters and improve computational efficiency.

4. (10 points) Which one of the statement is false about dropout?
- A. Dropout is an important regularization technique for training neural networks.
  - B. For neural networks that use dropout, the forward pass is slightly different between the training phase and inference phase.
  - C. During the inference phase, a neural network with dropout is equivalent to a smaller neural network with some neurons removed from computation.
  - D. Dropout layers in PyTorch contain a probability (hyper-)parameter as an input argument.
- 

Solution:

5. (10 points) Which statement is false about normalization and initialization.
- A. Appropriate initialization and normalization help alleviate the exploding/vanishing gradient issue.
  - B. If we don't use batch normalization, the training dynamics may suffer from slow convergence due to oscillation.
  - C. Xavier initialization and He initialization ensures that gradients have similar magnitude during the backward pass.
  - D. Batch normalization contains trainable parameters, which are updated using the usual backpropagation.
- 

Solution:

6. (10 points) Suppose that you use a CNN to process a minibatch of images. The batch size is 32. One convolutional layer is created as follows.

```
conv_layer = torch.nn.Conv2d(in_channels=4, out_channels=2,  
                           kernel_size = (5, 5), stride = (2, 2))
```

What is the **total** number of parameters in this layer?

---

Solution:

7. (10 points) Suppose that  $\sigma$  is a generic differentiable activation function. Suppose

- $x$  is a  $d$ -dimensional input vector;
- $W$  is a matrix of size  $d \times d$ .

Consider a nonstandard neural network

$$f(x) = x + Wx + W\sigma(Wx).$$

The activation functions are computed in a coordinate-wise way. We compute the loss based on MSE:

$$\ell(W, w) = (y - f(x)^\top w)^2.$$

We are interested in the gradient of  $\ell(W, w)$  with respect to  $W$ . Can you derive the gradient? Please do one of the two following tasks: (1) use the chain rule to write down the gradient calculation, (2) or draw a computational graph to represent the forward pass from  $x$  to the loss value.

---

Solution:

8. (10 points) Suppose that we have a neural network model `model` in PyTorch with three layers named “conv\_1”, “conv\_2”, “fc”. We know that we can access the layers by typing `model.conv_1`, `model.conv_2`, and `model.fc` respectively.

Suppose that the model is already trained, and we want to finetune the “fc” layer for a downstream application. To accomplish this, we need to freeze the parameters of the other layers before finetuning. How do you achieve this in PyTorch? Please write code below. *Hint: use `requires_grad`.*

---

Solution:

9. (20 points) Please explain in words what the following PyTorch code means.
- (i) `torchvision.transforms.Resize(size=(32,32))`
  - (ii) `class MLP(torch.nn.Module):`
  - (iii) `loss.backward()`
  - (iv) `optimizer = torch.optim.AdamW(model.parameters(), weight_decay=0.01)`
  - (v) `torch.manual_seed(20)`
- 

Solution:

