



STAT 453: Introduction to Deep Learning and Generative Models

Ben Lengerich

Lecture 18: Diffusion Models

November 5, 2025

Reading: See course homepage



Project

- <https://adaptinfer.github.io/dgm-fall-2025/project/>




Today

- Diffusion Models

Generative and Discriminative Models

- **Generative:**
 - Models the joint distribution $P(X, Y)$.
- **Discriminative:**
 - Models the conditional distribution $P(Y|X)$.






Where we're going: Deep Generative Models



Discriminative Model (what we've seen so far)



Generative Model (what we're going to see)



Gemini



Grok




deepseek



Modern Deep Generative Models (DGMs)

- Goal: Generative models of the form $P(X, Y, \theta)$ without strong simplifying assumptions.
- **Hidden structure z that explains high-dim. x**
- Fundamental challenge: We never observe z
- This makes two core computations difficult:
 - **Marginal likelihood:** $p_\theta(x) = \int p_\theta(x, z) dz$
 - **Posterior inference:** $p_\theta(z | x) \propto p_\theta(x | z)p(z)$
- Each type of DGM makes a tradeoff

Autoencoders




[Michelucci 2022]


Variational Autoencoders

Kullback-Leibler divergence term
where $p(z) = \mathcal{N}(\mu = 0, \sigma^2 = 1)$

$$L^{[i]} = -\mathbb{E}_{z \sim q_w(z|x^{[i]})} [\log p_w(x^{[i]}|z)] + \text{KL}(q_w(z|x^{[i]}) \| p(z))$$



Generative Adversarial Networks



Discriminator: $\max_D \mathcal{L}_D = \mathbb{E}_{\mathbf{x} \sim p_{data}(\mathbf{x})} [\log D(\mathbf{x})] + \mathbb{E}_{\mathbf{x} \sim G(\mathbf{z}), \mathbf{z} \sim p(\mathbf{z})} [\log(1 - D(\mathbf{x}))]$

Generator: $\min_G \mathcal{L}_G = \mathbb{E}_{\mathbf{x} \sim G(\mathbf{z}), \mathbf{z} \sim p(\mathbf{z})} [\log(1 - D(\mathbf{x}))]$.




Summary

Property	VAE	GAN
What we specify	Prior $p(z)$, Likelihood $p_\theta(x z)$	Prior $p(z)$, Generator $G_\theta(z)$
Induced $p(x)$	$p_\theta(x) = \int_z p_\theta(X z)p(z)dz$	$p_\theta(x) = \int_z p_\epsilon(x - G_\theta(z))p(z)dz$
Simplifying assumption	Choose a restricted variational posterior $q_\phi(z x)$	Replace NLL with a distributional discrepancy on samples (adversarial/IPM).
Training objective	ELBO: $E_q[\log p_\theta(x z)] - KL(q_\phi(z x) p(z))$	Minimax fooling discriminator
What's ignored from $p_\theta(x)$	$KL(q_\phi(z x) p_\theta(z x))$	All of NLL: $\log p_\theta(x)$ isn't evaluated or maximized.
Modes	Covering	Collapse
Generated Samples	Blurry	Realistic
Training	Relatively robust	Fragile



Diffusion Models

Overview and comparison of generative models



Diffusion



Diffusion models: forward pass

Turn a clear image into complete random noise

条件概率分布 (给定上一刻的图, 这一刻的图会变成啥样)

Normal distribution (Random noise)


$$q(\mathbf{x}_t | \mathbf{x}_{t-1}) = \mathcal{N}(\mathbf{x}_t; \sqrt{1 - \beta_t} \mathbf{x}_{t-1}, \beta_t \mathbf{I})$$

\downarrow
 Mean part
 (keep most of information from \mathbf{x}_{t-1})

For whole procedure of $\mathbf{x}_0 \rightarrow \mathbf{x}_T$

$$q(\mathbf{x}_{1:T} | \mathbf{x}_0) = \prod_{t=1}^T q(\mathbf{x}_t | \mathbf{x}_{t-1})$$

Multiply by T times (Markov chain)




Diffusion models: reverse pass

$$p_{\theta}(\mathbf{x}_{0:T}) = p(\mathbf{x}_T) \prod_{t=1}^T p_{\theta}(\mathbf{x}_{t-1} | \mathbf{x}_t)$$

While generation chain
 from \mathbf{x}_1 to \mathbf{x}_0


Mech ↑
 Variance ↑

Assume also follow the Normal distribution in backward PDI




Diffusion models: generating a new sample

$$p_{\theta}(\mathbf{x}_{0:T}) = p(\mathbf{x}_T) \prod_{t=1}^T p_{\theta}(\mathbf{x}_{t-1}|\mathbf{x}_t) \quad p_{\theta}(\mathbf{x}_{t-1}|\mathbf{x}_t) = \mathcal{N}(\mathbf{x}_{t-1}; \boldsymbol{\mu}_{\theta}(\mathbf{x}_t, t), \boldsymbol{\Sigma}_{\theta}(\mathbf{x}_t, t))$$




Should we really run this process on pixels?



★ The previous pixel diffusion is too hard to compute for large images
 → (Latent diffusion combine the AE with diffusion model)

Stable Diffusion: Add Text Conditioning



Latent Diffusion

Previously, we have unconditional generation, you cannot what it generates

Now, with text, it takes into account on how to describe the picture and generate conditionally



Stable Diffusion: Modern Image Generators





More reading

<https://lilianweng.github.io/posts/2021-07-11-diffusion-models/>

<https://lilianweng.github.io/posts/2021-07-11-diffusion-models/>

<https://theaisummer.com/diffusion-models/>



Property	VAE	GAN	Diffusion
What we specify	Prior $p(z)$, Likelihood $p_\theta(x z)$	Prior $p(z)$, Generator $G_\theta(z)$	Fixed forward noising $q(x_t x_{\{t-1\}})$; learn reverse $p_\theta(x_{t-1} x_t)$
Induced $p(x)$	$p_\theta(x) = \int_z p_\theta(X z) p(z) dz$	$p_\theta(x) = \int_z p_\epsilon(x - G_\theta(z)) p(z) dz$	$p_\theta(x) = \int p(x_T) \prod_t p_\theta(x_{t-1} x_t) dx$
Simplifying assumption	Choose a restricted variational posterior $q_\phi(z x)$	Replace NLL with a distributional discrepancy on samples (adversarial/IPM).	Fix forward noise q ; and optimize a variational bound on $-\log p_\theta(x_0)$.
Training objective	ELBO: $E_q[\log p_\theta(x z)] - KL(q_\phi(z x) \ p(z))$	Minimax fooling discriminator	VLB / score matching: with Gaussian schedules reduces to $\mathbb{E}_{t,x_0,\epsilon}[w(t) \ \epsilon - \epsilon_\theta(x_t, t) \ ^2]$
What's ignored from $p_\theta(x)$	$KL(q_\phi(z x) \ p_\theta(z x))$	All of NLL: $\log p_\theta(x)$ isn't evaluated or maximized.	Exact NLL not computed; optimize a variational upper bound on NLL (equivalently lower bound on $\log p$;(practical losses often reweight or drop constants from the exact VLB.
Modes	Covering	Collapse	Covering

Questions?

